

DD2460 – Graded Lab Assignment

Memory Safety

1 Overview

In this assignment, you will analyze *chibicc*, a tiny C compiler.

Description: *chibicc* is a program that converts source code written in the C language into assembly language, that is, a C compiler. The compiler itself is also developed using C.

URL: <https://github.com/cesarsotovalero/chibicc>

Note: Please do not download the source from the official repository, as it does not contain the injected fault, so there would be not much to see.

1.1 Installation and example

1. Unpack the sources:

```
$ git clone https://github.com/cesarsotovalero/chibicc
$ cd chibicc
```

2. Compile the sources:

```
$ make
```

3. Run the tool for the *nqueens* example:

```
$ chmod +x chibicc
$ ./chibicc examples/nqueens.c > tmp.s
```

4. (Later:) Compile the generated intermediate code:

```
$ gcc -static -o tmp tmp.s
```

5. (Later:) Run the example:

```
$ chmod +x tmp
$ ./tmp
```

Note: In the given version, *chibicc* will crash and leave an empty file. The last two steps will only run successfully on Linux *after* you have solved the assignment. They are not required for this exercise. You can also execute the script `run.sh`, which executes all the previous tasks.

2 Valgrind usage

Install valgrind (either as a package for your OS, or by downloading it from <https://valgrind.org/>).

Run Valgrind on the binary with the *nqueens* example as input:

```
$ valgrind --leak-check=full ./chibicc examples/nqueens.c > tmp.s
```

Valgrind will help you to understand why the *chibicc* crashes. You will see an error trace similar to this one:

```
==95176== Memcheck, a memory error detector
==95176== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==95176== Using Valgrind-3.14.0.GIT and LibVEX; rerun with -h for copyright info
==95176== Command: ./chibicc examples/nqueens.c
==95176==
--95176-- run: /usr/bin/dsymutil "./chibicc"
==95176== valgrind: Unrecognised instruction at address 0x1002cf4db.
==95176== at 0x1002CF4DB: __abort (in /usr/lib/system/libsystem_c.dylib)
==95176== by 0x1002CFD7D: __stack_chk_fail (in /usr/lib/system/libsystem_c.dylib)
==95176== by 0x10000940B: read_string_literal (tokenize.c:0)
==95176== by 0x10000902B: tokenize (tokenize.c:302)
==95176== by 0x1003052FF: ??? (in /dev/ttys007)
==95176== by 0x428: ???
==95176== Your program just tried to execute an instruction that Valgrind
==95176== did not recognise. There are two possible reasons for this.
==95176== 1. Your program has a bug and erroneously jumped to a non-code
==95176== location. If you are running Memcheck and you just saw a
==95176== warning about a bad jump, it's probably your program's fault.
==95176== 2. The instruction is legitimate but Valgrind doesn't handle it,
==95176== i.e. it's Valgrind's fault. If you think this is the case or
==95176== you are not sure, please let us know and we'll try to fix it.
==95176== Either way, Valgrind will now raise a SIGILL signal which will
==95176== probably kill your program.
==95176==
==95176== Process terminating with default action of signal 4 (SIGILL)
==95176== Illegal opcode at address 0x1002CF4DB
==95176== at 0x1002CF4DB: __abort (in /usr/lib/system/libsystem_c.dylib)
==95176== by 0x1002CFD7D: __stack_chk_fail (in /usr/lib/system/libsystem_c.dylib)
==95176== by 0x10000940B: read_string_literal (tokenize.c:0)
==95176== by 0x10000902B: tokenize (tokenize.c:302)
==95176== by 0x1003052FF: ??? (in /dev/ttys007)
==95176== by 0x428: ???
```

3 Explanation

There is an operation violating memory safety (in a serious way) inside `read_string_literal`.

Open the file `tokenize.c` and go to `read_string_literal` (line 180).

```
static Token *read_string_literal(Token *cur, char *start) {
    char *p = start + 1;
    char buf[32];
    int len = 0;

    for (;;) {
        if (*p == '\\0')
            error_at(start, "unclosed string literal");
        if (*p == '"')
            break;

        if (*p == '\\') {
            p++;
            buf[len++] = get_escape_char(*p++);
        } else {
            buf[len++] = *p++;
        }
    }
    ...
}
```

You can perhaps guess the meaning of the code of the `for` loop: It tries to read a string literal that has started with a double quote (") and exits when the closing quote is found. The loop also handles the backslash to escape the character following it, and copies over the resulting literal into the buffer `buf`. The buffer has a fixed length and is allocated on the stack.

4 Mandatory task: Error analysis

1. Explain the memory layout of the stack-allocated variables, the frame pointer, and the return pointer, in `read_string_literal`. Use the lecture slides from the memory safety lecture as a guide.
2. Explain what happened and why `valgrind` stops the program here.
3. Why is this potentially dangerous when running in an environment without memory safety checking?

5 Optional tasks

5.1 Fix (+0.5 points)

Observe that the size of `buf`, given by `sizeof(buf)`, is fixed. Access is controlled by `len`.

Deliverable: Write a patch that fixes the memory safety issue, using `error_at` to abort parsing when necessary.

5.2 Fuzzing (+0.5 points)

1. Clone `radamsa` from its home page: <https://gitlab.com/akihe/radamsa>

```
git clone https://gitlab.com/akihe/radamsa
cd radamsa
make
```

You do not need to run `make install`.

2. Set the environment variable `RADAMSA` so it points to the executable. Example:

```
export RADAMSA=$PWD/bin/radamsa
```

3. Change back to the location where you have `chibicc`.
4. Create a variant of the example input program *without* the long string:

```
grep -v Memory examples/nqueens.c > examples/base.c
```

5. Use fuzzing on the base example and inspect the result and/or run `valgrind` as above. You can create the example as follows:

```
$RADAMSA examples/base.c > examples/fuzzed.c
```

Try this step ten times and look into what types of perturbations you get.

Deliverable: Comment on your results:

1. What types of changes does `radamsa` introduce?
2. How many times has a string literal been changed?
3. Are the odds of getting a string literal that results in a buffer overflow any good?